Stochastic Control
Sequential Learning Algorithms
Course Information

D Manjunath & Jayakrishnan Nair

EE, IIT Bombay

January 4, 2022

Course information

- **Instructors:**
  D. Manjunath and Jayakrishnan Nair
  Office: EE 237-C, EE 237-D
  email: `dmanju@ee.iitb.ac.in, jayakrishnan.nair,`

Course information

- **Instructors:**
  D. Manjunath and Jayakrishnan Nair
  Office: EE 237-C, EE 237-D
  email: `dmanju@ee.iitb.ac.in, jayakrishnan.nair,`
- Teaching Assistants:

Course information

- **Instructors:**
  D. Manjunath and Jayakrishnan Nair
  Office: EE 237-C, EE 237-D
  email: `dmanju@ee.iitb.ac.in, jayakrishnan.nair,`
- Teaching Assistants:
  - ?

Course information

- **Instructors:**
  D. Manjunath and Jayakrishnan Nair
  Office: EE 237-C, EE 237-D
  email: `dmanju@ee.iitb.ac.in, jayakrishnan.nair,`
- Teaching Assistants:
  - ?
  - ??

Course information

- **Instructors:**
  D. Manjunath and Jayakrishnan Nair
  Office: EE 237-C, EE 237-D
  email: `dmanju@ee.iitb.ac.in, jayakrishnan.nair,`
- Teaching Assistants:
  - ?
  - ??
- **Course content**

Course information

- **Instructors:**
  D. Manjunath and Jayakrishnan Nair
  Office: EE 237-C, EE 237-D
  email: `dmanju@ee.iitb.ac.in, jayakrishnan.nair,`
- Teaching Assistants:
  - ?
  - ??
- **Course content**
  - Course will introduce theory of sequential learning, i.e., learning when data is presented in a sequence rather than in 'one-shot.'

Course information

- **Instructors:**
  D. Manjunath and Jayakrishnan Nair
  Office: EE 237-C, EE 237-D
  email: `dmanju@ee.iitb.ac.in, jayakrishnan.nair,`
- Teaching Assistants:
  - ?
  - ??
- **Course content**
  - Course will introduce theory of sequential learning, i.e., learning when data is presented in a sequence rather than in 'one-shot.'
  - Many of these algorihms are studied as **multi-armed bandit algorithms.**

Course information

- **Instructors:**
  D. Manjunath and Jayakrishnan Nair
  Office: EE 237-C, EE 237-D
  email: `dmanju@ee.iitb.ac.in, jayakrishnan.nair,`
- Teaching Assistants:
  - ?
  - ??
- **Course content**
  - Course will introduce theory of sequential learning, i.e., learning when data is presented in a sequence rather than in 'one-shot.'
  - Many of these algorihms are studied as **multi-armed bandit algorithms.**
  - Content will be predominantly mathematical with emphasis on provable properties of algorithms and working through the proofs.

Course information

- **Instructors:**
  D. Manjunath and Jayakrishnan Nair
  Office: EE 237-C, EE 237-D
  email: `dmanju@ee.iitb.ac.in, jayakrishnan.nair,`
- Teaching Assistants:
  - ?
  - ??
- **Course content**
  - Course will introduce theory of sequential learning, i.e., learning when data is presented in a sequence rather than in 'one-shot.'
  - Many of these algorihms are studied as **multi-armed bandit algorithms.**
  - Content will be predominantly mathematical with emphasis on provable properties of algorithms and working through the proofs.
  - A strong background in probability theory will be assumed.

Course information

- **Evaluation**

Attendance and class participation	5%
Homework and Quizzes	20%
Midsem Exam	25%
Endsem/Project	50%

Course information

- **Evaluation**

Attendance and class participation	5%
Homework and Quizzes	20%
Midsem Exam	25%
Endsem/Project	50%

- There may be minor changes to these weights as the course progresses; they will be advertised sufficiently in advance.

Course information

- **Lecture and exam schedule**

Course information

- **Lecture and exam schedule**
    - Lectures will be on MSTeams whenever they are online.

Course information

- **Lecture and exam schedule**
  - Lectures will be on MSTeams whenever they are online.
  - We would like to have classroom lectures; will start on this as soon as the Institite permits. We will make every attempt to stream the lectures online.

Course information

- **Lecture and exam schedule**
  - Lectures will be on MSTeams whenever they are online.
  - We would like to have classroom lectures; will start on this as soon as the Institite permits. We will make every attempt to stream the lectures online.
  - The exams will be during the mid-semester and end-semester (if applicable) exam weeks and will be as scheduled by the Institute

Course information

- **Lecture and exam schedule**
  - Lectures will be on MSTeams whenever they are online.
  - We would like to have classroom lectures; will start on this as soon as the Institite permits. We will make every attempt to stream the lectures online.
  - The exams will be during the mid-semester and end-semester (if applicable) exam weeks and will be as scheduled by the Institute
  - If we choose the project option, they will have to be done in groups of at most two. You are expected to bring your own ideas to a project. Project ideas should be pitched to us during the week of 01 Feb.

Course information

- **Lecture and exam schedule**
  - Lectures will be on MSTeams whenever they are online.
  - We would like to have classroom lectures; will start on this as soon as the Institite permits. We will make every attempt to stream the lectures online.
  - The exams will be during the mid-semester and end-semester (if applicable) exam weeks and will be as scheduled by the Institute
  - If we choose the project option, they will have to be done in groups of at most two. You are expected to bring your own ideas to a project. Project ideas should be pitched to us during the week of 01 Feb.
    - A schedule for this will be made available soon

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:
    1. Perfect experts,

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:
  1. Perfect experts,
  2. Imperfect experts

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:
  1. Perfect experts,
  2. Imperfect experts
- Bandit setting: Rather than all experts simultaneously, we get to sample only one expert at any time and we need to home in the best. An adversarial analysis

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:
  1. Perfect experts,
  2. Imperfect experts
- Bandit setting: Rather than all experts simultaneously, we get to sample only one expert at any time and we need to home in the best. An adversarial analysis
  1. Adversarial bandits and the EXP3 algorithm

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:
  1. Perfect experts,
  2. Imperfect experts
- Bandit setting: Rather than all experts simultaneously, we get to sample only one expert at any time and we need to home in the best. An adversarial analysis
  1. Adversarial bandits and the EXP3 algorithm
- Stochastic Bandits

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:
    1. Perfect experts,
    2. Imperfect experts
- Bandit setting: Rather than all experts simultaneously, we get to sample only one expert at any time and we need to home in the best. An adversarial analysis
    1. Adversarial bandits and the EXP3 algorithm
- Stochastic Bandits
    1. Pure Exploration

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:
  1. Perfect experts,
  2. Imperfect experts
- Bandit setting: Rather than all experts simultaneously, we get to sample only one expert at any time and we need to home in the best. An adversarial analysis
  1. Adversarial bandits and the EXP3 algorithm
- Stochastic Bandits
  1. Pure Exploration
  2. Explore and Commit

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:
  1. Perfect experts,
  2. Imperfect experts
- Bandit setting: Rather than all experts simultaneously, we get to sample only one expert at any time and we need to home in the best. An adversarial analysis
  1. Adversarial bandits and the EXP3 algorithm
- Stochastic Bandits
  1. Pure Exploration
  2. Explore and Commit
  3. Explore and Exploit tradeoff

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:
  1. Perfect experts,
  2. Imperfect experts
- Bandit setting: Rather than all experts simultaneously, we get to sample only one expert at any time and we need to home in the best. An adversarial analysis
  1. Adversarial bandits and the EXP3 algorithm
- Stochastic Bandits
  1. Pure Exploration
  2. Explore and Commit
  3. Explore and Exploit tradeoff
  4. The UCB algorithm and variants

Course Topics

- Basic problem setting of online learning and performance measures for algoroithm.
- Learning from experts, adversarial analysis:
  1. Perfect experts,
  2. Imperfect experts
- Bandit setting: Rather than all experts simultaneously, we get to sample only one expert at any time and we need to home in the best. An adversarial analysis
  1. Adversarial bandits and the EXP3 algorithm
- Stochastic Bandits
  1. Pure Exploration
  2. Explore and Commit
  3. Explore and Exploit tradeoff
  4. The UCB algorithm and variants
  5. Lower bounds

Topics

- Extensions with structure among the arms

Topics

- Extensions with structure among the arms
    1. Contextual

Topics

- Extensions with structure among the arms
  1. Contextual
  2. Combinatorial

Topics

- Extensions with structure among the arms
  1. Contextual
  2. Combinatorial
- Bayesian Bandit

Topics

- Extensions with structure among the arms
  1. Contextual
  2. Combinatorial
- Bayesian Bandit
  1. ML, MAP, and Bayesian Estimation

Topics

- Extensions with structure among the arms
  1. Contextual
  2. Combinatorial
- Bayesian Bandit
  1. ML, MAP, and Bayesian Estimation
  2. Thompson sampling

Topics

- Extensions with structure among the arms
  1. Contextual
  2. Combinatorial
- Bayesian Bandit
  1. ML, MAP, and Bayesian Estimation
  2. Thompson sampling
  3. Rested Bandits as Bayesian Bandits and the Gittins Index

Topics

- Extensions with structure among the arms
  1. Contextual
  2. Combinatorial
- Bayesian Bandit
  1. ML, MAP, and Bayesian Estimation
  2. Thompson sampling
  3. Rested Bandits as Bayesian Bandits and the Gittins Index
- Markov Decisions Processes

Topics

- Extensions with structure among the arms
  1. Contextual
  2. Combinatorial
- Bayesian Bandit
  1. ML, MAP, and Bayesian Estimation
  2. Thompson sampling
  3. Rested Bandits as Bayesian Bandits and the Gittins Index
- Markov Decisions Processes
- Markov Bandits

Topics

- Extensions with structure among the arms
  1. Contextual
  2. Combinatorial
- Bayesian Bandit
  1. ML, MAP, and Bayesian Estimation
  2. Thompson sampling
  3. Rested Bandits as Bayesian Bandits and the Gittins Index
- Markov Decisions Processes
- Markov Bandits
  1. Rested bandits as Markov bandits

Topics

- Extensions with structure among the arms
  1. Contextual
  2. Combinatorial
- Bayesian Bandit
  1. ML, MAP, and Bayesian Estimation
  2. Thompson sampling
  3. Rested Bandits as Bayesian Bandits and the Gittins Index
- Markov Decisions Processes
- Markov Bandits
  1. Rested bandits as Markov bandits
  2. Restless bandits and dynamic optimisation

Topics

- Extensions with structure among the arms
  1. Contextual
  2. Combinatorial
- Bayesian Bandit
  1. ML, MAP, and Bayesian Estimation
  2. Thompson sampling
  3. Rested Bandits as Bayesian Bandits and the Gittins Index
- Markov Decisions Processes
- Markov Bandits
  1. Rested bandits as Markov bandits
  2. Restless bandits and dynamic optimisation
- Reinforcement Learning

Sources

- A significant part of the course will follow
  T. Lattimore and C. Szepesvári, “Bandit Algorithms,” available at
  `http://tor-lattimore.com/downloads/book/book.pdf`.

Sources

- A significant part of the course will follow
  T. Lattimore and C. Szepesvári, “Bandit Algorithms,” available at
  `http://tor-lattimore.com/downloads/book/book.pdf`.
- Some of the others will based on papers and online material that we will cite whenever we borrow material from them.

Sources

- A significant part of the course will follow
  T. Lattimore and C. Szepesvári, "Bandit Algorithms," available at
  `http://tor-lattimore.com/downloads/book/book.pdf`.
- Some of the others will based on papers and online material that we will cite whenever we borrow material from them.
- It is critical that you read from textbook and not just depend on the notes from the lectures.

More about the project

- Start researching for your projects immediately.

More about the project

- Start researching for your projects immediately.
- The topics mentioned in a preceding slide can give you starting points.

More about the project

- Start researching for your projects immediately.
- The topics mentioned in a preceding slide can give you starting points.
- Last minute preparations will invariably be a hash and it is best avoided.

More about the project

- Start researching for your projects immediately.
- The topics mentioned in a preceding slide can give you starting points.
- Last minute preparations will invariably be a hash and it is best avoided.
- It would be good to put in the extra effort and not tag on some work you are doing elsewehere or its minor variant.